

Jinho Park

AI Application Engineer

📍 Seoul, South Korea
✉ jinhoparkseoul@gmail.com
🌐 jinhopark8345.github.io
📁 jinho-park-9a511a1b9
🔗 jinhopark8345

AI Application Engineer focused on **Information Retrieval** and **scalable backend systems** — production-grade vector search, Agentic RAG, latency optimization. Strong fast-prototyping track record with **Prefect, Kubernetes, GitOps** and high-availability **FastAPI** services. Regularly collaborate with **international teams and clients** across timezones.

EXPERIENCE

AI Application Engineer · Boeing

Jan 2024 – Present

ML Team · Seoul, South Korea

- Architected a high-precision **Agentic RAG** for aviation regulatory compliance: dedicated **Qdrant** retrieval service + generation app, with a **hybrid sparse + dense** search pipeline outperforming a dense-only configuration on a 200-query MRR eval.
- Tuned Qdrant indexing and payload filtering to hold **p95 retrieval latency under 1000 ms on CPU-only infrastructure**; sustained embedding throughput with **BentoML** dynamic batching + **ONNX Runtime** + quantization.
- Shipped a **GitLab MR bot**, an **MCP server** exposing internal tools to LLM agents, a **few-shot (v)LLM PDF parser** for RAG ingestion, and **gitlab-model-registry** for versioning AI artifacts on GitLab.
- Implemented **ArgoCD + OpenShift GitOps** (deploy cycles 1h to 10min) and stood up **Prometheus/Grafana** + a Postgres-backed KPI tracker (CTR, query quality, latency SLOs).

Machine Learning Engineer · Lomin

May 2022 – Aug 2023

ML Team · Seoul, South Korea

- Architected a high-throughput serving layer on **NVIDIA Triton** (KV-caching + quantization) that accelerated Transformer inference by **20%**.
- Built a multi-modal **Table Extraction** system for financial documents, improving **F1 by 15%**, and shipped a **Document Understanding** framework with graph-based augmentation pipelines.
- Delivered ML solutions and PoCs for enterprise clients: Kyobo Life, Lina Life, Shinhan Asset Trust, KIPO, Samsung Fire & Marine, Hanwha Life, Samsung Life, KB Insurance.

Software Engineer Intern — OpenVINO Korea · Intel

Aug 2021 – Feb 2022

OpenVINO Korea · Seoul, South Korea

- Analyzed deep-learning model performance on Intel hardware; fixed **5+ major** issues in OpenVINO for Intel integrated GPUs.

EDUCATION

B.S., Electrical and Electronics Engineering · Konkuk University, Korea

Mar 2015 – Feb 2022

Thesis: *Further Optimize MobileNetV2 with Channel-wise Squeeze and Excitation*. CS coursework: Algorithms, Data Structures, OS, Embedded Systems, AI.

Exchange Student, Computer Science · University of Agder, Norway

Aug 2017 – Jan 2018

OPEN SOURCE

- **Add BROS model to Transformers** · huggingface/transformers#23190 — ported architecture, tokenization, and spatial pre-processing for document understanding.
- **OpenVINO GPU fix: prepare_output** · openvinotoolkit/openvino#7783 — correctness fix on the Intel-integrated-GPU inference path.

SKILLS

Languages Python · C++

Databases PostgreSQL · Qdrant · Redis

ML / Inference PyTorch · LangChain · Transformers · NVIDIA Triton · OpenVINO · ONNX Runtime · Quantization · MLflow

Backend & App FastAPI · BentoML · Streamlit · React

Infra & MLOps Prefect · Kubernetes · Docker · Helm · GitLab CI/CD · ArgoCD · OpenShift · Prometheus · Grafana