

Jinho Park

Machine Learning Engineer

Relocating to Norway — spouse is a Norwegian citizen

📍 Seoul, South Korea
✉ jinhoparkseoul@gmail.com
🌐 jinhopark8345.github.io
📄 jinho-park-9a511a1b9
🔗 jinhopark8345

Machine Learning Engineer shipping production AI systems end-to-end: **Agentic RAG**, **document processing pipelines**, and **MCP-based agent tooling**, with the MLOps to keep them running. Comfortable running iterative PoCs with clients and translating fuzzy asks into concrete application logic. Exchange semester at the University of Agder, Norway; moving back permanently. On the side, runs **InitNode**, a studio shipping AI systems from PoC to production.

EXPERIENCE

AI Application Engineer · Boeing

Jan 2024 – Present

AI Team · Seoul, South Korea

- **Agentic RAG for regulated-industry document compliance:** Architected and implemented a high-precision retrieval system on Qdrant, exposed as an MCP server with the generation app as MCP client, with a hybrid sparse + dense pipeline that outperformed dense-only on a 200-query MRR eval.
- **Low-latency CPU retrieval:** Tuned Qdrant indexing and payload filtering to hold p95 retrieval latency under 1000 ms on CPU-only infrastructure; sustained embedding throughput with BentoML dynamic batching + ONNX Runtime + quantization.
- **Internal ML/LLM tooling:** Shipped an internal code-review LLM bot, an MCP server exposing in-house tools to LLM agents, a few-shot (v)LLM PDF parser for RAG ingestion, and an internal ML artifact registry on GitLab.
- **GitOps and observability:** Cut deploy cycles from 1h to 10min with ArgoCD on OpenShift; instrumented the stack with Prometheus, Grafana, and a Postgres-backed KPI tracker for CTR, query quality, adoption, and latency.

Machine Learning Engineer · Lomin

May 2022 – Aug 2023

ML Team · Seoul, South Korea

- **End-to-end document processing pipeline (on-premise):** Architected the production pipeline that processed **10,000 documents/hour** at **p95 end-to-end latency under 1 second** for enterprise financial and government clients.
- **Inference + extraction:** Triton with KV-caching and quantization for a **20% speedup**; multi-modal Table Extraction at **+15% F1**; Document Understanding framework with graph-based augmentation.
- **Clients:** Kyobo Life, Lina Life, Shinhan Asset Trust, KIPO (production); Samsung Life, Hanwha Life, Samsung Fire & Marine, KB Insurance (PoCs).

Software Engineer Intern · Intel

Aug 2021 – Feb 2022

OpenVINO · Seoul, South Korea

- **OpenVINO GPU fixes:** Analyzed deep-learning model performance on Intel hardware and fixed 5+ major issues in OpenVINO for Intel integrated GPUs.

EDUCATION

B.S., Electrical and Electronics Engineering · Konkuk University, Korea

Mar 2015 – Feb 2022

Thesis: *Further Optimize MobileNetV2 with Channel-wise Squeeze and Excitation*. CS coursework: Algorithms, Data Structures, OS, Embedded Systems, AI. Degree span includes mandatory Korean military service.

Exchange Student, Computer Science · University of Agder, Norway

Aug 2017 – Jan 2018

OPEN SOURCE

Add BROS · Huggingface / Transformers: ported architecture, tokenization, and spatial pre-processing for document understanding.

prepare_output · OpenVINO: correctness fix on the Intel-integrated-GPU inference path.

CERTIFICATIONS & AWARDS

CKA: Certified Kubernetes Administrator · Linux Foundation, Dec 2024.

Kaggle Competition Silver Medal — UBC-OCEAN · ovarian cancer subtype classification, 28th of 1,724 teams, Jan 2024.

SKILLS

Languages Python · C++ · TypeScript

Databases PostgreSQL · Redis

Backend FastAPI

Frontend React · Next.js · Streamlit

Infra & CI/CD Kubernetes (Openshift) · Helm · Docker ·

ArgoCD · GitLab CI/CD · Prefect

Monitoring Prometheus · Grafana · Langfuse

ML PyTorch · LangChain · Transformers

Inference BentoML · Triton · OpenVINO · ONNX

Runtime · Quantization

LLM API Anthropic · OpenAI

Vector DB Qdrant

Spoken Korean (native) · English (fluent) · Norwegian (beginner)