

Jinho Park

AI Application Engineer

📍 Seoul, South Korea
✉ jinhoparkseoul@gmail.com
🌐 jinhopark8345.github.io
📁 jinho-park-9a511a1b9
🔗 jinhopark8345

AI Application Engineer focused on **Information Retrieval** and **scalable backend systems**, building production-grade vector search, Agentic RAG, and latency optimization. Validate ideas fast with working PoCs, then take them to production. Day-to-day stack: **Prefect**, **Kubernetes**, **GitOps**, and high-availability **FastAPI** services. Regularly collaborate with **international teams and clients** across timezones.

EXPERIENCE

AI Application Engineer · Boeing

Jan 2024 – Present

AI Team · Seoul, South Korea

- **Agentic RAG for aviation regulatory compliance:** Architected a high-precision retrieval system on Qdrant: dedicated retrieval service + generation app, with a hybrid sparse + dense pipeline that outperformed dense-only on a 200-query MRR eval.
- **Low-latency CPU retrieval:** Tuned Qdrant indexing and payload filtering to hold p95 retrieval latency under 1000 ms on CPU-only infrastructure; sustained embedding throughput with BentoML dynamic batching + ONNX Runtime + quantization.
- **Internal LLM tooling:** Shipped a GitLab MR bot, an MCP server exposing internal tools to LLM agents, a few-shot (v)LLM PDF parser for RAG ingestion, and gitlab-model-registry for versioning AI artifacts on GitLab.
- **GitOps and observability:** Implemented ArgoCD + OpenShift GitOps (deploy cycles 1h to 10min) and stood up Prometheus/Grafana plus a Postgres-backed KPI tracker for CTR, query quality, and latency SLOs.

Machine Learning Engineer · Lomin

May 2022 – Aug 2023

ML Team · Seoul, South Korea

- **High-throughput Transformer inference:** Architected a serving layer on NVIDIA Triton with KV-caching and quantization that accelerated inference by 20%.
- **Multi-modal Table Extraction:** Built a system for financial documents, improving F1 by 15%, and shipped a Document Understanding framework with graph-based augmentation pipelines.
- **Enterprise delivery:** Led ML solutions and PoCs for Kyobo Life, Lina Life, Shinhan Asset Trust, KIPO, Samsung Fire & Marine, Hanwha Life, Samsung Life, and KB Insurance.

Software Engineer Intern · Intel

Aug 2021 – Feb 2022

OpenVINO · Seoul, South Korea

- **OpenVINO GPU fixes:** Analyzed deep-learning model performance on Intel hardware and fixed 5+ major issues in OpenVINO for Intel integrated GPUs.

EDUCATION

B.S., Electrical and Electronics Engineering · Konkuk University, Korea

Mar 2015 – Feb 2022

Thesis: *Further Optimize MobileNetV2 with Channel-wise Squeeze and Excitation*. CS coursework: Algorithms, Data Structures, OS, Embedded Systems, AI.

Exchange Student, Computer Science · University of Agder, Norway

Aug 2017 – Jan 2018

OPEN SOURCE

Add BROS · Huggingface / Transformers

Ported architecture, tokenization, and spatial pre-processing for document understanding.

prepare_output · OpenVINO

Correctness fix on the Intel-integrated-GPU inference path.

SKILLS

Languages Python · C++ · TypeScript

Databases PostgreSQL · Qdrant · Redis

ML / Inference PyTorch · LangChain · Transformers · NVIDIA Triton · OpenVINO · ONNX Runtime · Quantization · MLflow

Backend & App FastAPI · BentoML · Streamlit · React

Infra & MLOps Prefect · Kubernetes · Docker · Helm · GitLab CI/CD · ArgoCD · OpenShift · Prometheus · Grafana